

# Biostat 203B Homework 5

Due Mar 20 @ 11:59PM

AUTHOR

Ningke Zhang 705834790

## Predicting ICU duration

Using the ICU cohort `mimiciv_icu_cohort.rds` you built in Homework 4, develop at least three machine learning approaches (logistic regression with enet regularization, random forest, boosting, SVM, MLP, etc) plus a model stacking approach for predicting whether a patient's ICU stay will be longer than 2 days. You should use the `los_long` variable as the outcome. Your algorithms can use patient demographic information (gender, age at ICU `intime`, marital status, race), ICU admission information (first care unit), the last lab measurements before the ICU stay, and first vital measurements during ICU stay as features. You are welcome to use any feature engineering techniques you think are appropriate; but make sure to not use features that are not available at an ICU stay's `intime`. For instance, `last_careunit` cannot be used in your algorithms.

1. Data preprocessing and feature engineering.

```
# load libraries
library(stacks)
library(tidymodels)
```

— Attaching packages — tidymodels 1.3.0 —

✓ broom	1.0.7	✓ recipes	1.1.1
✓ dials	1.4.0	✓ rsample	1.2.1
✓ dplyr	1.1.4	✓ tibble	3.2.1
✓ ggplot2	3.5.1	✓ tidyr	1.3.1
✓ infer	1.0.7	✓ tune	1.3.0
✓ modeldata	1.4.0	✓ workflows	1.2.0
✓ parsnip	1.3.1	✓ workflowsets	1.1.0
✓ purrr	1.0.4	✓ yardstick	1.3.2

— Conflicts — tidymodels\_conflicts() —

```
* purrr::discard() masks scales::discard()
* dplyr::filter() masks stats::filter()
* dplyr::lag() masks stats::lag()
* recipes::step() masks stats::step()
```

```
library(dplyr)
library(recipes)
library(workflows)
library(tune)
library(glmnet)
```

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-8

```
library(vip)
```

Attaching package: 'vip'

The following object is masked from 'package:utils':

vi

```
library(ranger)
library(future)
library(xgboost)
```

Attaching package: 'xgboost'

The following object is masked from 'package:dplyr':

slice

```
# read data
mimiciv_icu_cohort <- readRDS("../hw4/mimiciv_shiny/mimic_icu_cohort.rds") |>
  select(-c(intime,
            outtime,
            admittime,
            disctime,
            deathtime,
            admit_provider_id,
            edregtime,
            edouttime,
            anchor_age,
            anchor_year,
            anchor_year_group,
            last_careunit,
            discharge_location,
            hospital_expire_flag,
            dod,
            los))
```

```
) |>
mutate(los_long = as.factor(los_long)) |>
print(width = Inf)
```

# A tibble: 94,458 × 26

	subject_id	hadm_id	stay_id	first_careunit
	<int>	<int>	<int>	<fct>
1	10000032	29079034	39553978	Medical Intensive Care Unit (MICU)
2	10000690	25860671	37081114	Medical Intensive Care Unit (MICU)
3	10000980	26913865	39765666	Medical Intensive Care Unit (MICU)
4	10001217	24597018	37067082	Surgical Intensive Care Unit (SICU)
5	10001217	27703517	34592300	Surgical Intensive Care Unit (SICU)
6	10001725	25563031	31205490	Medical/Surgical Intensive Care Unit (MICU/SICU)
7	10001843	26133978	39698942	Medical/Surgical Intensive Care Unit (MICU/SICU)
8	10001884	26184834	37510196	Medical Intensive Care Unit (MICU)
9	10002013	23581541	39060235	Cardiac Vascular Intensive Care Unit (CVICU)
10	10002114	27793700	34672098	Other

	admission_type	admission_location	insurance	language
	<fct>	<fct>	<chr>	<chr>
1	EW EMER.	EMERGENCY ROOM	Medicaid	English
2	EW EMER.	EMERGENCY ROOM	Medicare	English
3	EW EMER.	EMERGENCY ROOM	Medicare	English
4	EW EMER.	EMERGENCY ROOM	Private	Other
5	Other	PHYSICIAN REFERRAL	Private	Other
6	EW EMER.	Other	Private	English
7	URGENT	TRANSFER FROM HOSPITAL	Medicare	English
8	OBSERVATION ADMIT	EMERGENCY ROOM	Medicare	English
9	SURGICAL SAME DAY ADMISSION	PHYSICIAN REFERRAL	Medicare	English
10	OBSERVATION ADMIT	PHYSICIAN REFERRAL	Medicaid	English

	marital_status	race	gender	intime_age	hematocrit	bicarbonate	wbc
	<chr>	<chr>	<chr>	<int>	<dbl>	<dbl>	<dbl>
1	WIDOWED	WHITE	F	52	41.1	25	6.9
2	WIDOWED	WHITE	F	86	36.1	26	7.1
3	MARRIED	BLACK	F	76	27.3	21	5.3
4	MARRIED	WHITE	F	55	38.1	22	15.7
5	MARRIED	WHITE	F	55	37.4	30	5.4
6	MARRIED	WHITE	F	46	NA	NA	NA
7	SINGLE	WHITE	M	76	31.4	28	10.4
8	MARRIED	BLACK	F	77	39.7	30	12.2
9	SINGLE	Other	F	57	34.9	24	7.2
10	<NA>	Other	M	56	34.3	18	16.8

	creatinine	chloride	sodium	potassium	glucose	respiratory_rate
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.7	95	126	6.7	102	24
2	1	100	137	4.8	85	27
3	2.3	109	144	3.9	89	24
4	0.6	108	142	4.2	112	18
5	0.5	104	142	4.1	87	17
6	NA	98	139	4.1	NA	19
7	1.3	97	138	3.9	131	17

	8	1.1	88	130	4.5	141	16
9	0.9	102	137	3.5	288	14	
10	3.1	NA	125	6.5	95	22	

	non_invasive_blood_pressure_diastolic	heart_rate	temperature_fahrenheit
	<dbl>	<dbl>	<dbl>
1	48	91	98.7
2	63	80	97.7
3	127	77	98
4	90	86	98.5
5	97	96	97.6
6	56	86	97.7
7	85	131	97.9
8	49	60	98.1
9	70	80	97.2
10	80	111	97.9

	non_invasive_blood_pressure_systolic	los_long
	<dbl>	<fct>
1	84	FALSE
2	107	TRUE
3	158	FALSE
4	151	FALSE
5	167	FALSE
6	73	FALSE
7	112	FALSE
8	180	TRUE
9	104	FALSE
10	112	TRUE

# i 94,448 more rows

- Partition data into 50% training set and 50% test set. Stratify partitioning according to `los_long`. For grading purpose, sort the data by `subject_id`, `hadm_id`, and `stay_id` and use the seed `203` for the initial data split. Below is the sample code.

```
set.seed(203)

mimiciv_icu_cohort <- mimiciv_icu_cohort |>
  arrange(subject_id, hadm_id, stay_id) |>
  select(-c(subject_id, hadm_id, stay_id))
mimiciv_icu_cohort <- mimiciv_icu_cohort |> drop_na()

data_split <- initial_split(mimiciv_icu_cohort,
                           strata = "los_long",
                           prop = 0.5)

icu_other <- training(data_split)
icu_test <- testing(data_split)
```

- Train and tune the models using the training set.

- For logistic regression with elasticnet regularization, has 0.574 accuracy and 0.605 AUC in train data, the most important features are non invasive blood pressure systolic, frist care unit, and heart rate.
- For random forest, has 0.596 accuracy and 0.635 AUC in train data, the most important features are creatinine, intime age, and non invasive blood pressure systolic.
- For boosting, has 0.601 accuracy and 0.638 AUC in train data, the most important features are non invasive blood pressure systolic, intime age, and hematocrit.

```
# read models
logit_mod <- readRDS("final_fit_logistic_lastfit.rds")
rf_mod <- readRDS("final_fit_rf_lastfit.rds")
gb_mod <- readRDS("final_fit_gb_lastfit.rds")

logit_metrics <- logit_mod |> collect_metrics() |>
  filter(.metric %in% c("roc_auc", "accuracy"))
rf_metrics <- rf_mod |> collect_metrics() |>
  filter(.metric %in% c("roc_auc", "accuracy"))
gb_metrics <- gb_mod |> collect_metrics() |>
  filter(.metric %in% c("roc_auc", "accuracy"))

print(logit_metrics)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>        <dbl> <chr>
1 accuracy binary        0.574 Preprocessor1_Model1
2 roc_auc  binary        0.605 Preprocessor1_Model1
```

```
print(rf_metrics)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>        <dbl> <chr>
1 accuracy binary        0.596 Preprocessor1_Model1
2 roc_auc  binary        0.635 Preprocessor1_Model1
```

```
print(gb_metrics)
```

```
# A tibble: 2 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>        <dbl> <chr>
1 accuracy binary        0.601 Preprocessor1_Model1
2 roc_auc  binary        0.638 Preprocessor1_Model1
```

```
### Stacking model
recipe <-
  recipe(los_long ~ ., data = icu_other) |>
  step_impute_median(all_numeric_predictors()) |>
  step_impute_mode(all_nominal_predictors()) |>
```

```

step_novel(all_nominal_predictors()) |>
step_unknown(all_nominal_predictors()) |>
step_dummy(all_nominal_predictors()) |>
step_nzv(all_predictors()) |>
step_normalize(all_numeric_predictors(), -all_outcomes())

folds <- vfold_cv(icu_other, v = 2, strata = los_long)

#Logistic regression with elasticnet regularization
logit_mod <- logistic_reg(penalty = tune(), mixture = tune()) |>
  set_engine("glmnet", standardize = FALSE) |>
  set_mode("classification")

logit_wf <- workflow() |>
  add_recipe(recipe) |>
  add_model(logit_mod)

param_grid <- grid_regular(
  penalty(range = c(-6, 2)),
  mixture(range = c(0, 1)),
  levels = 5
)

logit_stacked <- logit_wf |>
  tune_grid(
    resamples = folds,
    grid = param_grid,
    metrics = metric_set(roc_auc, accuracy),
    control = control_stack_grid()
  )
logit_stacked

#Random Forest
rf_mod <-
  rand_forest(
    mode = "classification",
    mtry = tune(),
    trees = tune(),
    min_n = tune()) |>
  set_engine("ranger", importance = "permutation")

rf_wf <- workflow() |>
  add_recipe(recipe) |>
  add_model(rf_mod)

rf_grid <- grid_regular(
  mtry(range = c(2, 6)),
  trees(range = c(150, 200)),
  min_n(range = c(5, 10)),
  levels = 3
)

```

```

rf_stacked <- rf_wf |>
  tune_grid(
    resamples = folds,
    grid = rf_grid,
    metrics = metric_set(roc_auc, accuracy),
    control = control_stack_grid()
  )
rf_stacked

#Boosting
gb_mod <-
  boost_tree(
    mode = "classification",
    trees = 600,
    tree_depth = tune(),
    learn_rate = tune()
  ) |>
  set_engine("xgboost")

gb_wf <- workflow() |>
  add_recipe(recipe) |>
  add_model(gb_mod)

gb_grid <- grid_regular(
  tree_depth(range = c(3L, 8L)),
  learn_rate(range = c(-3, -0.5), trans = log10_trans()),
  levels = 5
)

gb_stacked <- gb_wf |>
  tune_grid(
    resamples = folds,
    grid = gb_grid,
    metrics = metric_set(roc_auc, accuracy),
    control = control_stack_grid()
  )
gb_stacked

```

```

class(logit_stacked)
class(rf_stacked)
class(gb_stacked)

# define the stacking model
stacked_model <- stacks() |>
  add_candidates(logit_stacked) |>
  add_candidates(rf_stacked) |>
  add_candidates(gb_stacked)

stacked_model <- stacked_model |>

```

```

blend_predictions(penalty = 1e-4, metric = metric_set(roc_auc, accuracy)) |>
fit_members()

#plot the stacked model
autoplot(stacked_model)

# compute the performance of the stacked model
#auc
stacked_results_prob <- stacked_model |>
  predict(new_data = icu_test, type = "prob")
stacked_results_prob
stack_auc <- stacked_results_prob |>
  bind_cols(icu_test) |>
  roc_auc(truth = los_long, .pred_TRUE, event_level = "second")
stack_auc
#accuracy
stacked_results_class <- stacked_model |>
  predict(new_data = icu_test, type = "class")
stack_acc <- stacked_results_class |>
  bind_cols(icu_test) |>
  accuracy(truth = los_long, estimate = .pred_class)
stack_acc

```

4. Compare model classification performance on the test set. Report both the area under ROC curve and accuracy for each machine learning algorithm and the model stacking. Interpret the results. What are the most important features in predicting long ICU stays? How do the models compare in terms of performance and interpretability?

```

# Summary of the performance

logit_auc <- logit_metrics |> filter(.metric == "roc_auc") |> pull(.estimate)
rf_auc <- rf_metrics |> filter(.metric == "roc_auc") |> pull(.estimate)
gb_auc <- gb_metrics |> filter(.metric == "roc_auc") |> pull(.estimate)
stack_auc <- stack_auc$.estimate

logit_acc <- logit_metrics |> filter(.metric == "accuracy") |> pull(.estimate)
rf_acc <- rf_metrics |> filter(.metric == "accuracy") |> pull(.estimate)
gb_acc <- gb_metrics |> filter(.metric == "accuracy") |> pull(.estimate)
stack_acc <- stack_acc$.estimate

Models <- c("Logit_enet", "Random Forest", "XGBoost", "Stacked")
ROC_AUC <- c(logit_auc, rf_auc, gb_auc, stack_auc)
Accuracy <- c(logit_acc, rf_acc, gb_acc, stack_acc)


model_performance <- data.frame(Models, ROC_AUC, Accuracy) |>
  mutate(across(where(is.numeric), round, digits = 4)) |>
  print()

```



### Summary:

```
Warning: There was 1 warning in `mutate()` :-  
x In argument: `across(where(is.numeric), round, digits  
= 4)`.  
Caused by warning:  
x The `...` argument of `across()` is deprecated as of  
dplyr 1.1.0.  
Supply arguments directly to `...` via `...` through an  
anonymous function instead.
```



Description: 47 (X = 2)

Model	RDC_AUC	Accuracy
Logit	0.6052	0.5742
Random Forest	0.6347	0.5965
XGBoost	0.6277	0.6013
Boosted	0.6455	0.6046

1/1/2015

data.frame

4 x 3

Description: df [4 × 3]

Models <chr>	ROC_AUC <dbl>	Accuracy <dbl>
Logit_enet	0.6052	0.5742
Random Forest	0.6347	0.5965
XGBoost	0.6377	0.6010
Stacked	0.6433	0.6046

4 rows

According to the prediction results of the four models, Stacking achieved the best performance (AUC = 0.6433, Accuracy = 60.46%), indicating that combining multiple models effectively enhances prediction accuracy. Logistic Regression (AUC = 0.6052) had the lowest performance, suggesting that linear relationships alone may not be sufficient to capture the complex patterns of ICU stay duration. XGBoost (~0.638 AUC) slightly outperformed Random Forest (~0.635 AUC), implying that boosting-based approaches are more effective in this dataset. So our most important features will follow the vip results of XGBoost: non invasive blood pressure systolic, intime age, and hematocrit.